

VRF-Based Mining

Simple Non-Outsourceable Cryptocurrency Mining

(Last Update: June 17, 2020)

Runchao Han^{1,2}, Haoyu Lin^{3,4}, and Jiangshan Yu¹

¹ Monash University

² CSIRO-Data61

³ Bytom Foundation

⁴ KZen Research

Abstract. This paper introduces VRF-based mining, a simple and effective way of making pooled mining impossible. Instead of using hash functions, VRF-based mining uses Verifiable Random Functions (VRFs) for Proof-of-work (PoW)-based consensus. As VRF binds an output with a secret key, a pool operator should reveal his private key to outsource the mining process to miners, and miners can anonymously steal cryptocurrency in the pool operator’s wallet.

We revisit the definition of *non-outsourceability* in existing research, and informally prove that VRF-based mining achieves the strongest notion of *non-outsourceability*. In addition, we discuss several considerations on instantiating VRFs for VRF-based mining. Moreover, we experimentally show that VRF-based mining is simple to implement and introduces negligible overhead.

1 Introduction

Bitcoin [24] started the era of cryptocurrency. In Bitcoin, each participant maintains its own ledger. The ledger is structured as a blockchain, i.e., a chain of blocks of transactions. Participants keep executing Proof-of-Work (PoW)-based consensus to agree on blocks and append them to the blockchain. In PoW-based consensus, each participant keeps trying to solve a computationally hard PoW puzzle [13]. Once solving the puzzle, the participant appends his block to the blockchain and gets some coins as reward. By adaptively adjusting the difficulty of PoW puzzles, PoW-based consensus can ensure only one participant solves the puzzle for each time period with high probability. Participants of PoW-based consensus are also known as miners, the process of solving PoW puzzles is known as mining, and the computation power used for mining is known as mining power.

Due to the limited rate of producing blocks, a miner’s reward can be highly volatile if he chooses to mine by himself. To stabilise mining rewards, miners usually choose to join mining pools. A mining pool allows miners to mine on blockchain in the name of the pool operator, and the pool operator distributes the mining reward to solo miners in a fine-grained way. While miners are willing to participate in mining pools, mining pools lead to the centralisation of

mining power, which contradicts Bitcoin’s design objective - decentralisation. Mining power centralisation weakens PoW-based consensus’ security. A mining pool with sufficient mining power can perform various attacks to break PoW-based consensus, such as selfish mining [14] and 51% attacks [24]. In addition, mining pools can censor transactions, i.e., decide which transactions to include in the blockchain. Nowadays, four largest mining pools control more than 51% Bitcoin mining power [4], which can be a lurking threat to blockchain ecosystem.

This paper introduces *VRF-based mining*, a surprisingly simple but effective approach to make pooled mining impossible. VRF-based mining replaces hash functions in PoW-based consensus with Verifiable Random Functions (VRFs) [22]. VRF requires a miner to input its secret key to compute the output, so a pool operator should reveal his secret key to miners in order to outsource mining. Thus, any miner can steal mining reward of a block owned by the pool operator anonymously. Concretely, our contributions are as follows.

VRF-based mining. We replace the hash function in PoW-based consensus with a VRF. As mining requires secret keys, the pool operator should publish his secret key to outsource mining, which enables miners to steal his mining reward. Verifying the VRF hash can prove both the correctness of the hash and the authorship of the block in the same time.

Non-outsourcability analysis. Miller et al. [23] first formalises the *non-outsourcability* of mining in PoW-based consensus. We find *non-outsourcability* consists of two orthogonal properties, namely the punishment on pool operators and the unaccountability of punishing. We then informally prove VRF-based mining achieves the strongest guarantee on both properties.

Instantiating VRF-based mining. VRF has four customisable components, namely the elliptic curve, two hash functions between strings and elliptic curve elements, and a normal hash function. We discuss considerations on choosing them for VRF-based mining.

Feasibility of VRF-based mining. We implement Elliptic curve-based VRF specified in [15] using Go programming language, and compare its performance with three mainstream mining algorithms SHA256D, Scrypt and CryptoNight. Our results show that VRF-based mining is easy to implement and introduces little overhead.

Partial outsourcing is unrealistic. Partial outsourcing is that, a pool operator interactively outsources computation that does not need the secret key to miners. We show that partial outsourcing is unprofitable, as it’s both computation-intensive and I/O-intensive.

This paper is organised as follows. §2 and §3 describes preliminaries and the construction of VRF-based mining, respectively. §4 revisits the definition of non-outsourcability, and shows VRF-based mining achieves better non-outsourcability than existing proposals. §5 discusses concerns of instantiating VRFs for VRF-based mining. §6 provides an experimental analysis on the practicality of VRF-based mining. §7 discusses why partial outsourcing in VRF-based mining is unprofitable. §8 discusses potential problems in VRF-based mining and their solu-

tions. §9 summaries related work, and §10 concludes this paper. We attach the pseudocode of the standardised EC-VRF construction [15] in Appendix.

2 Preliminaries

2.1 PoW-based consensus and mining

Miners participate in PoW-based consensus by *mining*, which works as follows. Given the blockchain’s latest view, the miner creates a block template t , which consists of transactions to include, hash of the last block, and other metadata. The miner then tries to solve a PoW puzzle with t as input. To solve this PoW puzzle, the miner should find a nonce n such that $H(t||n)$ is smaller than a difficulty parameter T , where $H(\cdot)$ is a hash function specified by the PoW-based consensus. As $H(\cdot)$ is one-way and its output is pseudorandom, the miner can only try different nonces until finding a valid one. Once the miner find a valid nonce, he can assemble this nonce and t to a block, and append this block to the blockchain. The miner will get some coins as the reward of mining. By adaptively adjusting T , the rate of new blocks remains stable. If miners mine too many blocks in a previous time period, the protocol will increase T so that miners will mine slower. Otherwise, the protocol will decrease T .

2.2 Mining pools

As the rate of new blocks is limited, the mining reward of miners is highly volatile, especially for miners with less powerful hardware. To stabilise mining reward, miners usually participate in mining pools. Mining pools enable miners to mine collaboratively and distribute reward in a fine-grained way. A mining pool usually relies on a centralised pool operator. All miners mine in the name of the pool operator, and the pool operator distributes rewards in proportion to miners’ contribution.

Concretely, a mining pool works as follows. First, a pool operator specifies the pool difficulty PT — a difficulty lower than the blockchain network — and a search interval $[n_1, n_m]$ of nonces. Then, the pool operator sends the block template t , PT , n_1 and n_m to the miner, and the miner starts to find a nonce in $[n_1, n_m]$ which satisfies PT . Once finding a valid nonce n_k , the miner sends n_k back to the pool operator. The pool operator then verifies whether n_k produces a hash satisfying PT , and records the miner’s contribution (a.k.a. a share) if valid. After a time period (say 24 hours), the pool operator calculates the total contribution of each miner, and distributes the mining reward to miners according to their submitted shares. As PT is easier to solve, calculating the mining power of a miner using shares is more fine-grained than using blocks. In this way, each miner is rewarded in a more fine-grained way, so more stably.

2.3 Verifiable random functions

Verifiable Random Function (VRF) [22] is a public-key version of cryptographic hash function. In addition to the input string, VRF involves a pair of a secret

key and a public key. Given an input string and a secret key, one can compute a hash. Anyone knowing the associated public key can verify the correctness of the hash, and can also verify the hash is generated by the owner of the secret key. Formally, a VRF consists of four algorithms: VRFKeyGen , VRFHash , VRFProve and VRFVerify .

- $\text{VRFKeyGen}(1^\lambda) \rightarrow (sk, pk)$: On input security parameter 1^λ , outputs a secret/public key pair (sk, pk) .
- $\text{VRFHash}(sk, \alpha) \rightarrow \beta$: On input sk and an arbitrary-length string α , outputs a fixed-length hash β .
- $\text{VRFProve}(sk, \alpha) \rightarrow \pi$: On input sk and α , outputs the proof π for β .
- $\text{VRFVerify}(pk, \alpha, \beta, \pi) \rightarrow \{True, False\}$: On input pk , α , β , π , outputs the verification result *True* or *False*.

Informally, VRF should preserve the following three security properties [15]:

- *Uniqueness*: Given a secret key sk and an input α , $\text{VRFHash}(sk, \alpha)$ produces a unique valid output.
- *Collision Resistance*: It is computationally hard to find two inputs α and α' that $\text{VRFHash}(sk, \alpha) = \text{VRFHash}(sk, \alpha')$.
- *Pseudorandomness*: It is computationally hard to distinguish the output of $\text{VRFHash}(sk, \alpha)$ from a random string if not knowing the corresponding public key pk and proof π .

Algorithm 3 in Appendix A describes the Elliptic-curve-based VRF (EC-VRF) construction standardised in draft-goldbe-vrf [15].

3 VRF-based mining

Instead of using a hash function, VRF-based mining uses a VRF to produce hashes that satisfy the difficulty requirement. VRF takes both an input and a secret key to produce a hash. The owner of this secret key can produce a proof proving the ownership of his hash. Thus, to outsource the mining process, the pool operator has to reveal his secret key to miners. However, with the secret key, any miner in the mining pool can steal all mining rewards.

Cryptocurrency mining consists of two components, namely mining blocks and verifying blocks. We call the process of mining a block *Work*, and the process of verifying a block *Verify*. Algorithm 1 and 2 describe *Work* and *Verify* of VRF-based mining, respectively.

Work. Miners run *Work* to mine new blocks. More specifically, a miner - with his private key sk the block template (a complete block without nonce) t - keeps searching for a nonce n that can make the (VRF) hash *out* of the block *blk* to meet the blockchain difficulty T . Once finding a valid n , the miner produces the proof π of *out*, and appends *blk* (with n), *out* and π to the blockchain.

Algorithm 1: $\text{Work}(sk, t, T)$.

Input: The secret key sk , the block template t , and the blockchain difficulty T
Output: The block blk , the VRF output out , and the VRF proof π
Initialise n , out , blk ▷ Initialise variables
while $n+ = 1$ **do** ▷ Refresh the nonce
 $blk \leftarrow \text{ConstructBlock}(t, n)$ ▷ Assemble block
 $out \leftarrow \text{VRFHash}(sk, blk)$ ▷ Produce VRF output
 if $out < T$ **then** ▷ If meeting difficulty
 break ▷ Mining successful
 end
end
 $\pi \leftarrow \text{VRFPProve}(sk, blk)$ ▷ Produce the proof
return blk , out , π ▷ Return block, hash and proof

Algorithm 2: $\text{Verify}(blk, out, \pi, T)$

$pk \leftarrow blk.txs[0].scriptPubKey$ ▷ Find miner's pk
require($out < T$) ▷ Check if satisfying diff
/* Here $\text{VRFVerify}(\cdot)$ ensures: */
/* 1. out is generated by the owner of sk */
/* 2. out is a valid output of $\text{VRFHash}(sk, blk)$ */
require($\text{VRFVerify}(pk, blk, out, \pi)$)
... ▷ Verify other fields
... ▷ Verify transactions

Verify. Upon incoming blocks, miners run **Verify** to verify their validity. While other verifications are the same as in hash-based mining, **Verify** in VRF-based mining should additionally run $\text{VRFVerify}(\cdot)$ to verify 1) whether out is produced by the owner of sk , and 2) whether out is a valid output of $\text{VRFHash}(sk, blk)$.

Block structure. Different from hash-based mining, in VRF-based mining a block should additionally attach out and π , but no longer need the signature of a block. Other miners without knowing sk cannot produce out , but can use π to verify out is generated by someone knowing sk . Through proving the authorship of out , π also proves the authorship of the block. Thus, miners only need to sign out , but do not need to sign blocks.

4 Non-outsourcability analysis

In this section, we revisit existing definitions on *non-outsourcability*, and show that VRF-based mining achieves the strongest notion of *non-outsourcability*.

4.1 Revised definitions

Miller et al. [23] first formalise cryptocurrency mining as *scratch-off puzzles*, and formally define two levels of *non-outsourcability*, namely *weak non-outsourcability* and *strong non-outsourcability*.

- *Weak non-outsourcability*: If the pool operator outsources the mining process, miners can always steal the reward of mining.
- *Strong non-outsourcability*: In addition to *weak non-outsourcability*, the pool operator cannot link the stolen mining reward with the miner who steals it.

Basically, *weak non-outsourcability* defines the punishment of outsourcing, while *strong non-outsourcability* covers both the punishment and the anonymity of the stealer. We call the property defining the punishment of outsourcing *punish-mining-reward*. The anonymity of stealers defined in [23] is a special case of *transaction unlinkability* [29]: Given two arbitrary transactions, it is hard to know whether their outputs belong to the same secret key. To steal the mining reward in [23], the miner should create a transaction, of which the input is the mining reward and the output points to his address. Finding out who steals the mining reward is equivalent to finding out the address holding the output, and we call the property making such attempt infeasible *stealing-unlinkability*.

4.2 Non-outsourcability of VRF-based mining

VRF-based mining achieves both *punish-mining-reward* and *stealing-unlinkability*. In VRF-based mining, the pool operator outsources mining by revealing his secret key to miners. The best thing the pool operator can do is to use a new secret for each block, and stealers can only steal the mining reward of a single block with a secret key. Thus, VRF-based mining achieves *punish-mining-reward*. Similar with the Strongly non-outsourcable scratch-off puzzle [23], VRF-based mining achieves *Stealing-unlinkability* by allowing a stealer to use a freshly generated address to steal cryptocurrency. To receive the stolen cryptocurrency of the pool operator, the stealer can create a new address, and construct a transaction of which the stolen cryptocurrency directs to this address. As this new address has no historical transactions, linking the transaction stealing cryptocurrency with other transactions can be impossible. Then the stealer can then spend his stolen cryptocurrency anonymously using numerous techniques, such as mixing services [21][11][25][16] and stealth addresses [29].

5 Instantiating VRF

In order to implement VRF-based mining, one needs to first instantiate the VRF. VRF in Algorithm 3 has four configurable components, including the elliptic curve and three hash functions. Three hash functions are: $H_1(\cdot)$ mapping an arbitrary-length string into an elliptic curve element, $H_2(\cdot)$ mapping an elliptic curve element into a fixed-length string, and $H_3(\cdot)$ mapping an arbitrary-length string into a fixed-length string. We discuss considerations on choosing these four components for VRF-based mining.

5.1 Elliptic curve

As neither blockchains and VRF limits the choice of elliptic curves, any elliptic curve can be adapted. For example, among prominent elliptic curves, Ed25519 [8] can be a promising choice. First, Ed25519 works on a prime field with the prime number $2^{255} - 19$, which provides sufficient enumeration space for VRF. Second, Ed25519 supports EdDSA, a fast and secure digital signature algorithm. Last, Ed25519 is compatible with existing blockchains, as numerous blockchains and projects using VRF adapt Ed25519 as their underlying elliptic curve.

5.2 Hashing a string to an elliptic curve point $H_1(\cdot)$

The first hash function $H_1(\cdot)$ hashes a string to an elliptic curve point. It should prevent adversaries from learning any knowledge of the input from its hash. In addition, the $H_1(\cdot)$ function used in VRF should be deterministic, otherwise the hash will be unreproducible.

A standardisation document [26] specifies several hash-to-curve functions [17, ?, ?, ?] that fit into different elliptic curves and satisfy our requirements. Within these functions, Elligator2 [9] is the recommended one for Ed25519.

5.3 Hashing an elliptic curve point to a string $H_2(\cdot)$

The second hash function $H_2(\cdot)$ hashes an elliptic curve point to a fixed-length string. It can be divided to two steps: 1) encoding an elliptic curve point to a string, and 2) hashing the string using a normal hash function. The encoding step can be bidirectional and unencrypted, so can be done simply by converting a big integer to a string. The hashing step should be cryptographically secure, so can adapt any existing cryptographically secure hash functions.

5.4 Normal hash function $H_3(\cdot)$

The third hash function $H_3(\cdot)$ hashes an arbitrary string to a fixed-length string. It is only used in VRFProve and VRFVerify i.e., producing and verifying proofs of VRF outputs. The overhead of proving and verification should be minimised, so fast and cryptographically secure hash functions such as Keccak [10] and BLAKE [7] are suitable for $H_3(\cdot)$.

5.5 Memory-hard mining

Some PoW-based blockchains — such as Ethereum [33] and Monero [3] — employ memory-hard hash functions in order to minimise advantage of specialised mining hardware and achieve better decentralisation. Different from normal hash functions, the performance of computing memory-hard hash functions is bounded by memory access rather than arithmetic operations in processors. As memory access is much harder to optimise compared to arithmetic operations

in processors, high-end hardware cannot outperform low-end hardware greatly in terms of memory-hard functions.

To make VRF-based mining memory-hard, VRFHash should be memory-hard. VRFHash of the standardised VRF consists of one $H_1(\cdot)$ hashing, one scalar-point multiplication and one $H_2(\cdot)$ hashing. By making $H_1(\cdot)$ or $H_2(\cdot)$ memory-hard, VRF-based mining will be memory-hard. This can be achieved by embedding a memory-hard hash function such as Ethash [32] and CryptoNight [27] in $H_1(\cdot)$ or $H_2(\cdot)$.

6 Practicality of VRF-based mining

In this section, we experimentally show that VRF-based mining is easy to implement and introduces small overhead. First, we show that VRFs can be much more efficient than hash functions used for mining, including Scrypt in Litecoin and CryptoNight in Monero. Second, we profile the performance of VRFHash, and show that the elliptic curve scalar multiplication dominates VRFHash’s runtime, which can be optimised in the future.

6.1 Experimental setting

We implement the standardised EC-VRF in Algorithm 3 using Go programming language, without any optimisation. We use Ed25519 [8] as the underlying elliptic curve, Elligator [9] as $H_1(\cdot)$, SHA-3 as the hash function. Ed25519, SHA-3 and the encoding of points on Ed25519 are supported by Go’s standard library. We use open-source implementations of SHA256D ⁵, Scrypt ⁶, and CryptoNight ⁷. All experiments run on a MacBook Pro with a 2.2 GHz Intel Core i7 Processor and a 16 GB DDR4 RAM. Each group of experiments consists of ten runs, and we take the average value of ten values as the result.

6.2 VRF v.s. existing mining algorithms

We compare the performance of VRF with existing mining algorithms. Figure 1a shows the runtime of VRF and other algorithms. SHA256D is much faster than other algorithms, as SHA256D is simply executing SHA256 twice. In addition, VRFVerify is slightly slower than VRFProve, and VRFProve is slightly slower than VRFHash. Last, although without optimisation, all functions of VRF are much faster than Scrypt and CryptoNight. This means that VRF is easy to implement and introduces little overhead, so suitable for cryptocurrency mining.

⁵ <https://github.com/seehuhn/sha256d>

⁶ <https://github.com/elithrar/simple-scrypt>

⁷ <https://github.com/Equim-chan/cryptonight>

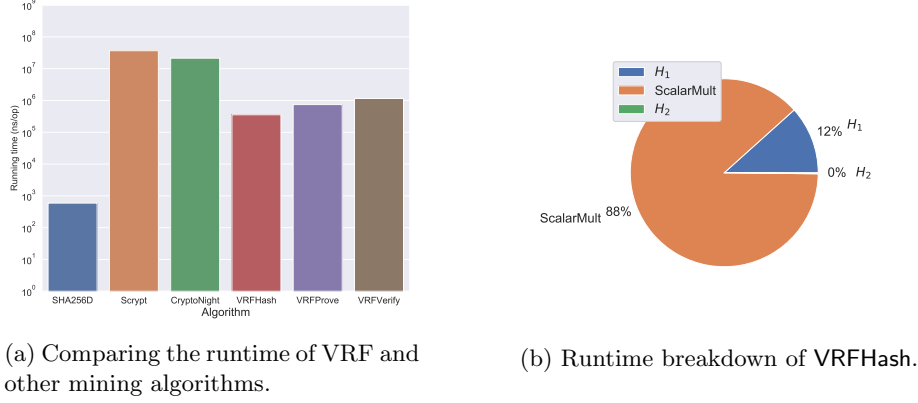


Fig. 1: Evaluation of VRF.

6.3 Runtime breakdown of VRF

We profile `VRFHash` by evaluating its runtime of each step. Figure 1b shows that, the elliptic curve scalar multiplication $\gamma \leftarrow h^{sk}$ takes 88% of `VRFHash`'s running time. This is because we use SHA-3 as the hash function in $H_1(\cdot)$ and $H_2(\cdot)$, and SHA-3 is designed to be fast. Meanwhile, we calculate the elliptic curve scalar multiplication using the trivial double-and-add method without any optimisation, thus is much slower than $H_1(\cdot)$ and $H_2(\cdot)$. Even so, `VRFHash` is still much faster than `Script` and `CryptoNight`. This further proves that VRFs introduce little overhead.

7 Profitability of partial outsourcing

The pool operator may still have opportunity to partially outsource mining. The mining function `VRFHash` consists of three steps: 1) $h \leftarrow H_1(\alpha)$, 2) $\gamma \leftarrow h^{sk}$ and 3) $\beta \leftarrow H_2(\gamma)$. *Non-outsourcability* is from the second step $\gamma = h^{sk}$, as it requires the knowledge of the pool operator's secret key sk . The pool operator can outsource the first step $h = H_1(\alpha)$ and the last step $\beta = H_2(\gamma)$ by distributing different α s and γ s to miners, respectively. However, we show that such partial outsourcing is very inefficient and unprofitable compared to outsourcing in hash-based mining, due to the computing and I/O overhead.

7.1 Partial outsourcing

In VRF-based mining, the pool operator can outsource $H_1(\cdot)$ or $H_2(\cdot)$. To outsource $H_1(\cdot)$, the pool operator generates a series of nonces $\{n_1, \dots, n_m\}$, then sends the series and the block template t to a miner. Then, the miner computes $H_1(n_i)$ for each $i \in [1, m]$, then sends back all $h_i = H_1(\cdot)$ hashes to the pool operator. The pool operator should verify $\{h_1, \dots, h_m\}$ before starting the second step — multiplying each of $\{h_1, \dots, h_m\}$ with his secret key sk .

After the second step, the pool operator obtains a series of $\{\gamma_1, \dots, \gamma_m\}$. To outsource $H_2(\cdot)$, the pool operator first sends $\{\gamma_1, \dots, \gamma_m\}$ as well as the pool difficulty PT to the miner. Then, the miner calculates $H_2(\gamma_i)$ for each $i \in [1, m]$, compare the hashes with PT , and sends back γ s that satisfy the difficulty (denoted as I). Upon receiving these γ s, the pool operator verifies their correctness and accumulates the mined shares to the miner's total contribution.

7.2 First obstacle: overhead of verification

The first obstacle of partial outsourcing is verifying hashes from miners. For outsourcing $H_1(\cdot)$, the pool operator should verify all of $\{h_1, \dots, h_m\}$. For outsourcing $H_2(\cdot)$, The pool operator should verify σ s satisfying PT . Such verification overhead is even more than the pool operator mining by himself. Thus, partial outsourcing unprofitable.

7.3 Second obstacle: overhead of I/O

To bypass the first obstacle, the pool operator and the miners can trust each other and omit the verification. This is possible as pooled mining is beneficial for them: the pool operator can earn some fees from the miners, while miners can stabilise their mining reward.

However, partial outsourcing can still be unprofitable even without the verification. VRF-based mining is extremely I/O intensive, and the network bandwidth of the pool operator's server is limited. To outsource a single $H_1(\cdot)$, the pool operator should at least receive a $H_1(\cdot)$ hash from the miner. To outsource a single $H_2(\cdot)$, the pool operator should at least send a σ to the miner.

Assume a pool operator with bandwidth BW (in bytes/s), and the pool operator can achieve optimal bandwidth utilisation on mining. Assume each the pool operator should transfer N bytes for outsourcing each mining attempt. Then, the maximum hashrate that the pool operator can achieve is

$$\text{Maximum hashrate} = \frac{BW}{N}$$

Either a $H_1(\cdot)$ hash or a γ (a point on the elliptic curve) is at least 32 bytes. If outsourcing both $H_1(\cdot)$ and $H_2(\cdot)$, the maximum hashrate the pool operator can support is $\text{Maximum Hashrate} = \frac{BW}{64}(h/s)$. If outsourcing only one of them, the maximum hashrate is $\text{Maximum Hashrate} = \frac{BW}{32}(h/s)$.

Figure 2 shows relationship between the server's network bandwidth and the maximum hashrate the server can achieve. The result shows that, existing servers can only achieve limited hashrate on partial outsourcing. On AWS [1], I3EN Metal is the server with most bandwidth, which is 3,500 MB/s. However, I3EN Metal can support hashrate less than $\frac{1}{10}$ of Monero's total mining power. Within those mainstream cryptocurrencies, Monero's hashrate is the least. Even if the pool operator rents a cluster of I3EN Metal machines for more bandwidth, I3EN Metal is quite expensive (10.848 USD per hour), leading to high expense of maintaining a mining pool. Therefore, partial outsourcing in VRF-based mining is costly and unrealistic.




Fig. 2: Server bandwidth v.s. maximum hashrate. Data of server bandwidth and hashrate were fetched from AWS [1] and CoinWarz [2] at 01/03/2020, respectively.

8 Discussions

While eliminating mining pools can improve decentralisation, it also introduces new problems. We discuss two problems and their possible solutions, namely 1) high reward variance and 2) secret key leakage in memory.

8.1 Weaker security guarantee of PoW-based consensus

The reason of miners joining mining pools is that, miners may not obtain stable income via solo mining. Miners may be discouraged to mine if their rewards are highly volatile. With weaker incentive of mining, fewer miners will join the blockchain. The blockchain will then attract less mining power, which weakens the security of PoW-based consensus.

This problem can be addressed by making the mining reward fine-grained. Miller et al. [23] proposes the *multi-tier reward scheme*. In multi-tier reward scheme, the mining difficulty is divided into different levels, and miners can mine blocks satisfying different difficulty levels arbitrarily. In this way, mining reward is distributed in a fine-grained way, so lowers the reward variance. StrongChain [28] introduces the notion of Collaborative PoW, where miners are encouraged to mine blocks together and the mining reward is distributed in proportion to miners' contributions.

Another solution is to increase the rate of mining blocks. With more blocks mined in a time unit, the mining reward can also be more stable. For example, although of independent interest, protocols for scaling blockchains such as sharding [31] and DAG [19] can also stabilise the mining reward variance.

Table 1: Comparison between mining protocols.

	VRF-based mining	Puzzle-1 [23]	Puzzle-2 [23]	2P-PoW [18]
Punishment	New reward	New reward	New reward	New reward
Stealing	Unlinkable	Linkable	Unlinkable	Unlinkable
No partial outsourcing	✓	✓	✓	✗
Support rand. signatures	✓	✓	✓	✗
No complex cryptography	✓	✓	✗	✓

8.2 Secret key leakage in memory

VRF-mining requires the secret key, so a miner should keep its secret key in plaintext in memory *all the time*. This gives adversaries opportunity to steal the secret key from the memory. For example, if the mining software has a bug that enables hackers to access the memory space of the mining software, then the hacker can easily steal the secret key.

For VRF-based mining, keeping secret keys in memory is inevitable. To achieve *non-outsourceability*, the secret key should be combined to the mining process. As long as mining does not stop, the secret key should be kept in plaintext in memory. This applies to all protocols that execute frequently and require secret keys, such as TLS.

There have been several different ways to protect in-memory secret keys. First, the miner can isolate the scalar multiplication to a software or hardware enclave. Note that the encalves are unnecessary to be general-purpose. Second, the process of the mining software can proactively destroy itself once detecting anomalous memory access attempts. Third, the miner can isolate mining from the blockchain node. For example, he can run the mining process independently with his blockchain node process. In this way, if an adversary compromises his node but not the operating system, the adversary still cannot read the private key in another process. Fourth, he can run the mining process in an offline machine that can only communicate with his blockchain node. Fifth, Oblivious RAM (ORAM) is a promising primitive to protect sensitive data in memory. ORAM allows CPUs to access data in memory while the data in memory is encrypted and the access pattern is hidden. Last, the miner can use a new secret key for each block, so that leaking a secret key only affects the reward of a single block.

9 Related work

To the best of our knowledge, VRF-based mining is the first construction that makes pooled mining *impossible*. We briefly review related research on preventing mining pools, and compare them with VRF-based mining. We classify related research to two types, namely pooled-mining-unfriendly mining protocols and decentralised mining pools.

Table 2: Comparison with decentralised mining pools.

	VRF-based mining	P2Pool [30]	SmartPool [20]	BetterHash [12]
Complexity	-	Blockchain	Smart contract	-
Decentralisation	Mining	Mining	Mining	Select txs

9.1 Mining protocols

There are two mining protocols aiming at discouraging or breaking mining pools: the *non-outsourceable scratch-off puzzle* [23] and the *Two Phase Proof-of-Work (2P-PoW)* [18]. Table 1 summarises our comparisons.

Non-outsourceable scratch-off puzzle. Miller et al. [23] formalises mining as *scratch-off puzzles*, defines *non-outsourceability*, and proposes two *non-outsourceable scratch-off puzzles*. One achieves *weak non-outsourceability* — miners can steal the mining reward, and the other achieves *strong non-outsourceability* — miners can steal the mining reward anonymously.

The weak non-outsourceable scratch-off puzzle works as follows. First, the miner randomly generates a Merkle tree. Second, the miner randomly chooses a nonce, samples some leaves of the Merkle tree according to the nonce, and hashes their Merkle proofs together to a single hash. If this hash meets the difficulty requirement, then the nonce is valid. Last, the miner binds the valid nonce and its block template together to a valid block.

In order to outsource the mining process, the mining pool should distribute the search space of nonces to miners. A miner can steal the mining reward by binding valid nonces it finds with his own block template. However, this stealing process is not anonymous. Unlike existing PoW-based consensus where miners can choose both nonces and block templates, the weak non-outsourceable scratch-off puzzle only allows miners to choose nonces, so all miners share the same search space of nonces. Thus, the pool operator can link the nonce in the stolen block with the miner who is assigned with the search space covering this nonce. To achieve *strong non-outsourceability*, the strong non-outsourceable scratch-off puzzle replaces the plaintext nonce in the block with a Zero Knowledge Proof proving the statement “I know a valid nonce”.

As discussed in §4, while being simpler and more efficient than non-outsourceable scratch-off puzzle, our VRF-based mining achieves *strong non-outsourceability*.

2P-PoW. Eyal and Sirer proposes 2P-PoW [18], a mining protocol that discourages pooled mining. In 2P-PoW, mining consists of two phases, and each phase has a difficulty parameter. A miner should find a nonce that makes the block to pass two phases: 1) the SHA256D hash of the block meets the first difficulty, 2) the SHA256 hash of the signature of the block meets the second difficulty. As the second requirement requires the private key, pool operators cannot outsource the second phase.

Compared to VRF-based mining that makes pooled mining impossible, 2P-PoW is partially outsourceable by outsourcing the first phase. In addition, 2P-PoW requires the digital signature algorithm to be deterministic, while commonly used digital signatures such as ECDSA rely on randomisation. If the

signature is randomised, the pool operator can make use of all nonces from miners that meet the first phase but fail the second phase. Given a nonce meeting the first difficulty, the pool operator repetitively generates signatures to meet the second requirement. Moreover, how to adjust two difficulties still remains unknown and requires some simulations.

9.2 Decentralised mining pools

Decentralised mining pool is a type of mining pools that work in a decentralised way: miners mine in the name of themselves and share reward in a fine-grained way. In this way, miners are rewarded stably while mining power is not controlled by pool operators. Table 2 summarises the comparison between VRF-based mining with decentralised mining pools.

P2Pool [30] is a decentralised mining pool for Bitcoin. Instead of using a centralised server, P2Pool uses a blockchain called *share-chain* to receive shares from miners. All miners in P2Pool participate in the PoW-based consensus of share-chain. Once a miner finds a share on Bitcoin, he appends a block to the share-chain. The block consists of the share and a coinbase transaction that rewards miners in proportion to their shares. Once a miner mines a block that also satisfies Bitcoin’s difficulty, he submits this block to the Bitcoin blockchain, and miners are rewarded according to the coinbase transaction. P2Pool suffers from several limitations. First, handling the difficulty of mining shares is hard. If the difficulty is high, miners’ reward will still be volatile. If the difficulty is low, there will be numerous low-difficulty shares, which introduces huge overhead on broadcasting and receiving shares. Second, shares on share-chain are much more frequent than blocks on Bitcoin blockchain, and the share-chain suffers from high orphan rate. Last, existing research shows that P2Pool is vulnerable to temporary dishonest majority [34].

SmartPool [20] is another decentralised mining pool. Instead of using a blockchain, SmartPool employs smart contracts to track shares from miners. This implies that SmartPool cannot work on blockchains without smart contracts. In addition, as blockchains achieve limited throughput, blockchains can only handle a limited number of shares for each time unit. In this way, distributing mining reward may also be delayed, especially when a large number of miners participate in the SmartPool. Moreover, as the SmartPool smart contract should verify the validity of blocks, miners should submit the entire block — including transactions — to the SmartPool. Verifying blocks introduces non-negligible computing overhead and transaction fee, and storing blocks in smart contracts also introduces non-negligible overhead on storage.

BetterHash [12] is another decentralised mining protocol, which has been integrated into *Stratum V2* [6], the next generation of the *Stratum* [5] pooled mining protocol. In *BetterHash*, the block operator allows miners to choose transactions and construct blocks in his name, rather than constructing block templates by himself. Thus, *BetterHash* only contributes to the decentralisation of constructing blocks, but does not improve the decentralisation of mining power.

10 Conclusion and future work

In this paper, we propose VRF-based mining, that can make pooled mining in Proof-of-work-based consensus impossible. VRF-based mining is simple and intuitive: miners produce hashes of blocks using VRFs rather than hash functions, so that a pool operator should reveal his secret key to outsource the mining process to other miners. We informally prove that VRF-based mining achieves the strongest notion of *non-outsourceability*. In addition, we discuss considerations of instantiating VRF-based mining. Moreover, we experimentally prove that VRF-based mining is simple to implement and introduces negligible overhead.

References

1. Amazon ebs-optimized instances, <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-optimized.html>
2. Cryptocurrency mining profitability calculator, <https://www.coinwarz.com/cryptocurrency>
3. Monero - secure, private, untraceable, <https://www.getmonero.org>
4. Pool stats - btc.com, <https://btc.com/stats/pool>, last accessed on 24/11/19
5. Stratum mining protocol - bitcoin wiki (2015), https://en.bitcoin.it/wiki/Stratum_mining_protocol, last accessed on 08/12/19
6. Stratum v2 — the next generation protocol for pooled mining (2019), <https://stratumprotocol.org/>, last accessed on 08/12/19
7. Aumasson, J.P., Henzen, L., Meier, W., Phan, R.C.W.: Sha-3 proposal blake. Submission to NIST **92** (2008)
8. Bernstein, D.J., Duif, N., Lange, T., Schwabe, P., Yang, B.Y.: High-speed high-security signatures. *Journal of Cryptographic Engineering* **2**(2), 77–89 (2012)
9. Bernstein, D.J., Hamburg, M., Krasnova, A., Lange, T.: Elligator: Elliptic-curve points indistinguishable from uniform random strings. In: *Proceedings of the 2013 ACM SIGSAC conference on Computer & communications security*. pp. 967–980. ACM (2013)
10. Bertoni, G., Daemen, J., Peeters, M., Van Assche, G.: Keccak. In: *Annual international conference on the theory and applications of cryptographic techniques*. pp. 313–314. Springer (2013)
11. Bonneau, J., Narayanan, A., Miller, A., Clark, J., Kroll, J.A., Felten, E.W.: Mixcoin: Anonymity for bitcoin with accountable mixes. In: *International Conference on Financial Cryptography and Data Security*. pp. 486–504. Springer (2014)
12. Corallo, M.: Betterhash mining protocol(s) (2019), <https://github.com/TheBlueMatt/bips/blob/betterhash/bip-XXXX.mediawiki>, last accessed on 08/12/19
13. Dwork, C., Naor, M.: Pricing via processing or combatting junk mail. In: *Annual International Cryptology Conference*. pp. 139–147. Springer (1992)
14. Eyal, I., Sirer, E.G.: Majority is not enough: Bitcoin mining is vulnerable. *Communications of the ACM* **61**(7), 95–102 (2018)
15. Goldberg, S., Papadopoulos, D., Vcelak, J.: draft-goldbe-vrf: Verifiable random functions.(2017) (2017)
16. Heilman, E., Alshenibr, L., Baldimtsi, F., Scafuro, A., Goldberg, S.: Tumblebit: An untrusted bitcoin-compatible anonymous payment hub. In: *Network and Distributed System Security Symposium* (2017)

17. Icart, T.: How to hash into elliptic curves. In: Annual International Cryptology Conference. pp. 303–316. Springer (2009)
18. Ittay, E., Sirer, Gün, E.: How to disincentivize large bitcoin mining pools (2014), <http://hackingdistributed.com/2014/06/18/how-to-disincentivize-large-bitcoin-mining-pools/>
19. Li, C., Li, P., Zhou, D., Xu, W., Long, F., Yao, A.: Scaling nakamoto consensus to thousands of transactions per second. arXiv preprint arXiv:1805.03870 (2018)
20. Luu, L., Velner, Y., Teutsch, J., Saxena, P.: Smartpool: Practical decentralized pooled mining. In: 26th {USENIX} Security Symposium ({USENIX} Security 17). pp. 1409–1426 (2017)
21. Maxwell, G.: Coinjoin: Bitcoin privacy for the real world. In: Post on Bitcoin forum (2013)
22. Micali, S., Rabin, M., Vadhan, S.: Verifiable random functions. In: 40th Annual Symposium on Foundations of Computer Science (Cat. No. 99CB37039). pp. 120–130. IEEE (1999)
23. Miller, A., Kosba, A., Katz, J., Shi, E.: Nonoutsourcable scratch-off puzzles to discourage bitcoin mining coalitions. In: Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security. pp. 680–691. ACM (2015)
24. Nakamoto, S., et al.: Bitcoin: A peer-to-peer electronic cash system (2008)
25. Ruffing, T., Moreno-Sanchez, P., Kate, A.: Coinshuffle: Practical decentralized coin mixing for bitcoin. In: European Symposium on Research in Computer Security. pp. 345–364. Springer (2014)
26. Scott, S., Sullivan, N., Wood, C.A.: Hashing to elliptic curves. Tech. rep., Internet-Draft draft-irtf-cfrg-hash-to-curve-03. Internet Engineering Task ... (2019)
27. Seigen, M.J., Nieminen, T.: Neocortex, and antonio m. juarez. cryptonight hash function. cryptonote, march 2013
28. Szalachowski, P., Reijsbergen, D., Homoliak, I., Sun, S.: Strongchain: Transparent and collaborative proof-of-work consensus. arXiv preprint arXiv:1905.09655 (2019)
29. Van Saberhagen, N.: Cryptonote v 2.0 (2013)
30. Voight, F.: p2pool: Decentralized, dos-resistant, hop-proof pool (2011)
31. Wang, J., Wang, H.: Monoxide: Scale out blockchains with asynchronous consensus zones. In: 16th {USENIX} Symposium on Networked Systems Design and Implementation ({NSDI} 19). pp. 95–112 (2019)
32. Wiki, E.: Ethash. GitHub Ethereum Wiki. <https://github.com/ethereum/wiki/wiki/Ethash> (2017)
33. Wood, G., et al.: Ethereum: A secure decentralised generalised transaction ledger. Ethereum project yellow paper **151**(2014), 1–32 (2014)
34. Zamyatin, A.: Decentralized mining pools: Security and attacks, https://www.alexeyzamyatin.me/files/Decentralized_Mining-Security_and_Attacks.pdf

A The standardised elliptic-curve-based VRF

Algorithm 3 describes the Elliptic-curve-based VRF (EC-VRF) construction standardised in draft-goldbe-vrf [15].

Algorithm 3: The Elliptic-curve-based VRF (EC-VRF) construction standardised in draft-goldbe-vrf [15].

Preliminaries:

- G is a cyclic group of prime order q with generator g .
- $H_1(\cdot)$ hashes an arbitrary-length string into an element in G .
- $H_2(\cdot)$ hashes an element in G into a fixed-length string.
- $H_3(\cdot)$ hashes an arbitrary-length string into a fixed-length string.
- $\text{random}([x, y])$ uniformly and randomly picks a number in $[x, y]$.

Algorithm $(sk, pk) \leftarrow \text{VRFKeyGen}(1^\lambda)$:

```

 $sk = \text{random}([0, q - 1])$ 
 $pk = g^{sk}$ 
return  $(sk, pk)$ 

```

Algorithm $\beta \leftarrow \text{VRFFHash}(sk, \alpha)$:

```

 $h = H_1(\alpha)$ 
 $\gamma = h^{sk}$ 
 $\beta = H_2(\gamma)$ 
return  $\beta$ 

```

Algorithm $\pi \leftarrow \text{VRFPProve}(sk, \alpha)$:

```

 $h = H_1(\alpha)$ 
 $\gamma = h^{sk}$ 
 $k = \text{random}([0, q - 1])$ 
 $c = H_3(g, h, g^{sk}, h^{sk}, g^k, h^k)$ 
 $s = k - c \cdot sk \pmod{q}$ 
 $\pi = (\gamma, c, s)$ 
return  $\pi$ 

```

Algorithm $\{0, 1\} \leftarrow \text{VRFVerify}(pk, \alpha, \beta, \pi)$:

```

 $u = pk^c \cdot g^s$ 
 $h = H_1(\alpha)$ 
if  $\gamma \notin G$  then
  | return 0
end
 $v = \gamma^c \cdot h^s$ 
if  $c \neq H_3(g, h, pk, \gamma, u, v)$  then
  | return 0
end
return 1

```
